

## Master on Foundations of Data Science



# Recommender Systems

Collaborative Recommender Systems (II)

Santi Seguí | 2019-2020

# Today

- RecSys Challenge Task
- How are we gonna survive to COVID-19
- Factorization Models

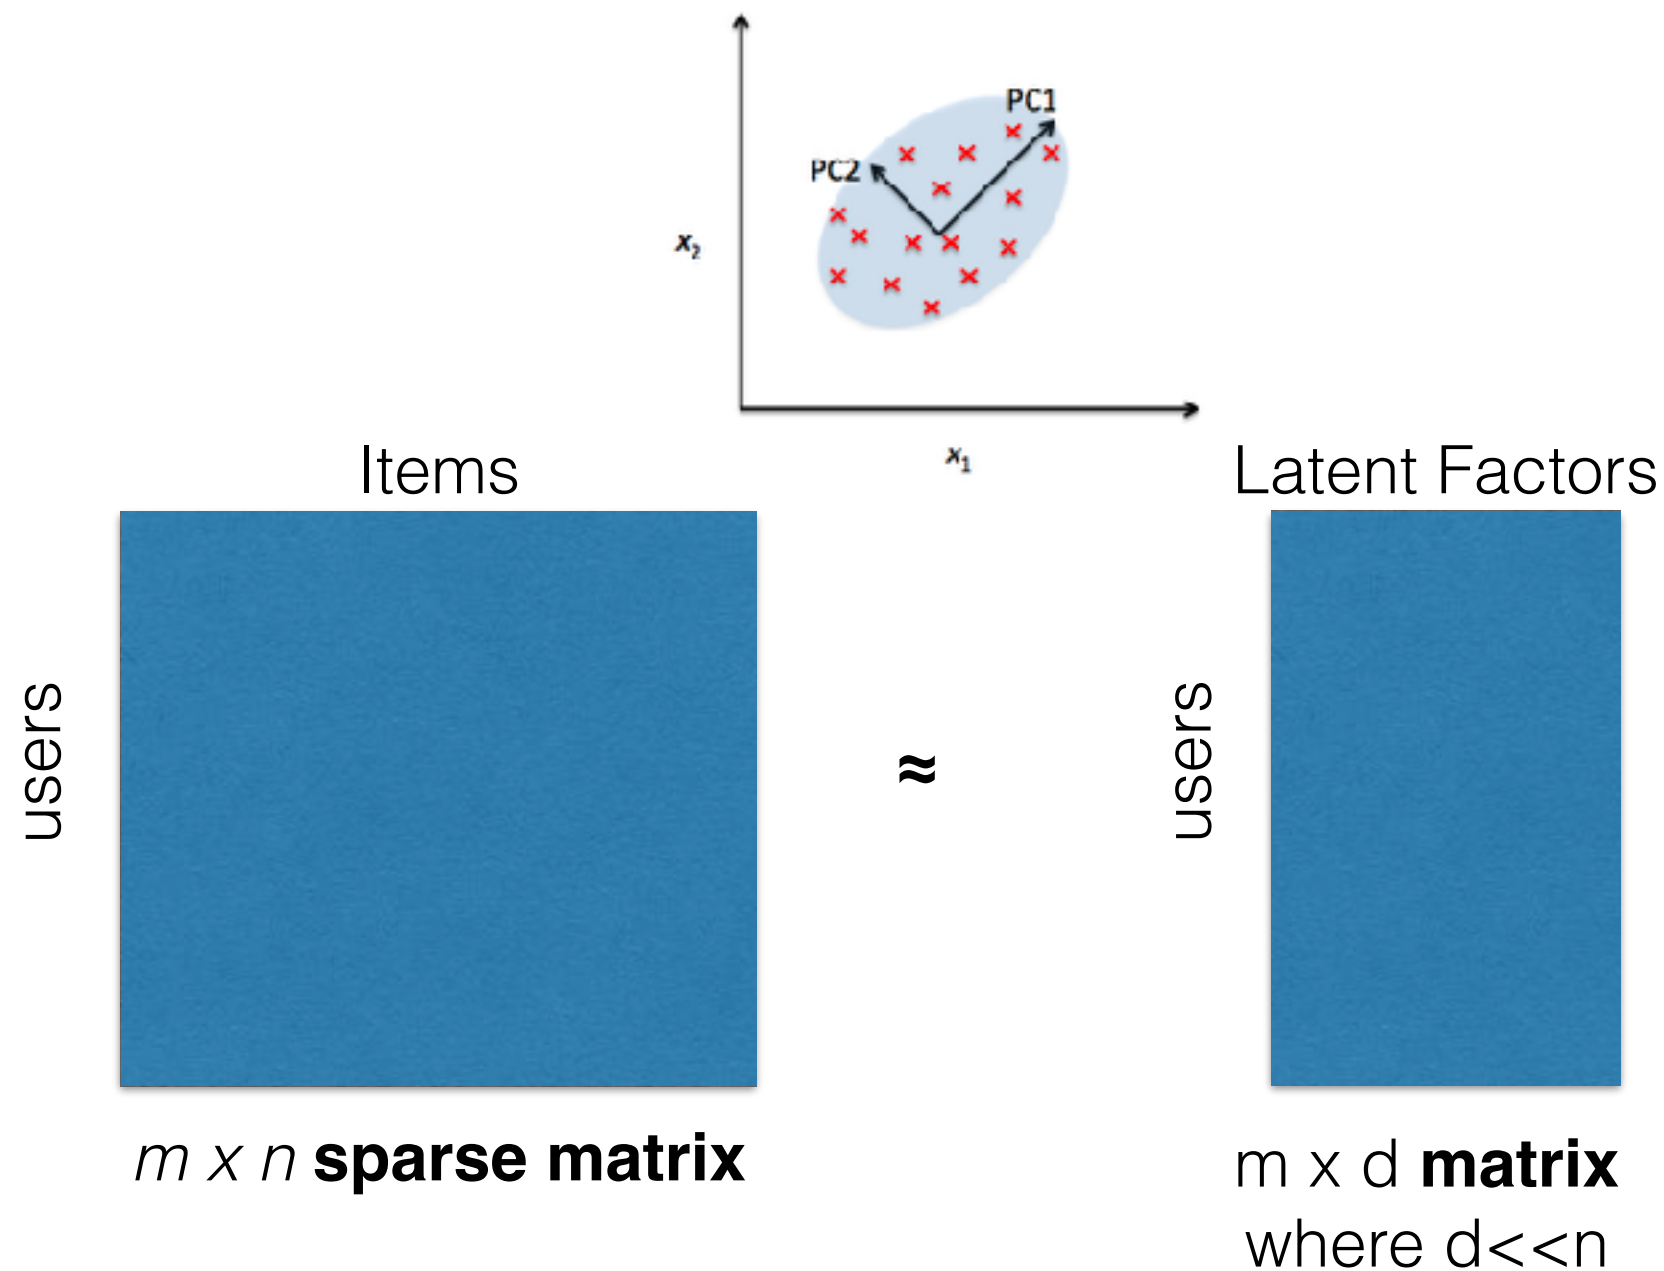
# How we will survive to COVID-19

- Course Content will be provided by:
  - Weekly call (20-30 min) + questions
  - Slides + Papers
  - Notebooks
- Evaluation will be based on:
  - RecSys Challenge + 2 new **individual** assignments
  - NO EXAM

# Dimensionality Reduction

- Pairwise similarities are hard to robustly be computed in sparse matrices.
- Dimensionality reduction can be used to **improve** neighborhood-based methods both in terms of **quality** and in terms of **efficiency**
- A reduced representation of the data can be created in terms of either row-wise latent factors or in terms of column-wise latent factors.

# Dimensionality Reduction



# Dimensionality Reduction

- The low-dimensional representation can be computed using **PCA** or **SVD-Like** methods.
- After the d-dimensional representation of each user is estimated, the similarity between users can be computed
- Cosine or dot product on the reduced vectors can be used in order to compute the similarity
- More robust since the feature vector is fully specified
- More efficient

# Dimensionality Reduction

- How to **obtain** the **d-dimensional representation** on the sparse matrix?
- **SVD Method**. Steps:
  - Augment the  $m \times n$  incomplete rating matrix  $R \rightarrow R_f$ 
    - Mean-user rating or mean-item rating for each row/column
  - Let's define the similarity matrix  $S$  as  $\mathbf{S} = \mathbf{R}_f^T \mathbf{R}_f$ .  $S$  is a positive semi-definite of size  $n \times n$
  - Determine the dominant basis vectors of  $R_f$  by computing the **diagonalization** of the similarity matrix  $S$ .
    - $S = P\Lambda P^T$ , where  $P$  is an  $n \times n$  matrix, whose columns contain the orthonormal eigenvectors of  $S$ .  $\Lambda$  is a diagonal matrix containing the non-negative eigenvalues of  $S$  along its diagonal.
  - Let denote  $P_d$  the  $n \times d$  matrix only containing the columns of  $P$  with the largest eigenvalues
  - The low representation of  $R$  is obtained by the multiplication of  $\mathbf{R}_f \mathbf{P}_d$

# Dimensionality Reduction

- How to **obtain** the **d-dimensional representation** on the sparse matrix?
- **PCA Method**. Steps:
  - Augment the  $m \times n$  incomplete rating matrix  $R \rightarrow R_f$ 
    - Mean-user rating or mean-item rating for each row/column
  - Let's define the similarity matrix  $S$  as **the Covariance Matrix of  $R_f$**
  - Determine the dominant basis vectors of  $R_f$  by computing the diagonalization of the similarity matrix  $S$ .
    - $S = P\Lambda P^T$ , where  $P$  is an  $n \times n$  matrix, whose columns contain the orthonormal eigenvectors of  $S$ .  $\Lambda$  is a diagonal matrix containing the non-negative eigenvalues of  $S$  along its diagonal.
  - Let denote  $P_d$  the  $n \times d$  matrix only containing the columns of  $P$  with the largest eigenvalues
  - The low representation of  $R$  is obtained by the multiplication of  **$R_f P_d$**



# Challenges on Factorization

- Challenges:
  - Missing Values
    - Need a way to fill it
    - Several alternatives, including clever averages and predictions
  - Computational Complexity
  - Lack of transparency/explainability


# Model-Based methods

# Model-Based methods

- **Neighborhood-based** models are very popular because its **simplicity**, today, these methods are not necessarily the most accurate ones.
  - based on distance between user/items.
  - Knn-generalization
- Today, model-based methods and in particular **latent factor models** are some of the most **accurate** methods
  - learn best parameters for the task

# Latent Factor Models

Figure 21.21



# Latent Factor Models

$$R = UV^T$$

$$R \approx UV^T$$

$$r_{ij} \approx \bar{u}_i \cdot \bar{v}_j$$

$$= \sum_{s=1}^k (\text{Affinity of user } i \text{ to concept } s) \cdot (\text{Affinity of item } j \text{ to the concept } s)$$

# Latent Factor Models

	Nero	Julius Cesar	Cleopatra	Sleepless in Seattle	Pretty Woman	Casablanca		History	Romance		Nero	Julius Cesar	Cleopatra	Sleepless in Seattle	Pretty Woman	Casablanca	
U1	1	1	1	0	0	0	=	1	0	<b>X</b>							
U2	1	1	1	0	0	0		1	0								
U3	1	1	1	0	0	0		1	0								
U4	1	1	1	1	1	1		1	1		History	1	1	1	0	0	0
U5	-1	-1	-1	1	1	1		-1	1		Romance	0	0	1	1	1	1
U6	-1	-1	1	1	1	1		-1	1								
U7	-1	-1	-1	1	1	1		-1	1								
	<b>R</b>							<b>U</b>			<b>V</b>						

# Model-Based methods

- Advantages over neighborhood methods:
  - **Space-efficiency.** Usually, the learned model is much smaller than the original rating matrix
  - **Training speed and prediction speed.** Neighborhood-based models takes quadratic complexity in either number of users or number of items
  - **Avoiding overfitting.** The summarized model use to help in avoiding overfitting. Regularization methods can be used to make these models robust

# Latent Factor Models

Matrix factorization models map both users and items to a joint latent factor space of dimensionality, such that user-item interactions are modeled as inner products in that space



# Matrix Factorization Methods

- Latent factor models approach tries to explain the ratings by characterizing both items and users to a small number of factors inferred from the rating patterns.
- Singular Value Decomposition (SVD) is a well established technique for identifying latent semantic factors. Done by factorizing the user-item rating matrix.
- **PROBLEM:** Hard to optimize due to the huge amount of missing values

# Singular Value Decomposition (SVD)

- Columns of  $U$  and  $V$  are constrained to be mutually orthogonal.
- Mutual orthogonality has the advantage that the concepts can be completely independent of one another. Can be interpreted in scatterplots

$$\begin{array}{c} \hat{X} \\ \left( \begin{array}{cccc} x_{11} & x_{12} & \dots & x_{1n} \\ x_{21} & x_{22} & \dots & \\ \vdots & \vdots & \ddots & \\ x_{m1} & & & x_{mn} \end{array} \right) \\ m \times n \end{array} \approx \begin{array}{c} U \\ \left( \begin{array}{ccc} u_{11} & \dots & u_{1r} \\ \vdots & \ddots & \\ u_{m1} & & u_{mr} \end{array} \right) \\ m \times r \end{array} \begin{array}{c} S \\ \left( \begin{array}{ccc} s_{11} & 0 & \dots \\ 0 & \ddots & \\ \vdots & & s_{rr} \end{array} \right) \\ r \times r \end{array} \begin{array}{c} V^T \\ \left( \begin{array}{ccc} v_{11} & \dots & v_{1n} \\ \vdots & \ddots & \\ v_{r1} & & v_{rn} \end{array} \right) \\ r \times n \end{array}$$

# Example of Singular Value Decomposition

$$R_f = \begin{pmatrix} 1 & -1 & 1 & -1 & 1 & -1 \\ 1 & 1 & ? & -1 & -1 & -1 \\ ? & 1 & 1 & -1 & -1 & ? \\ 1 & -1 & -1 & 1 & 1 & 1 \\ 1 & ? & -1 & 1 & 1 & 1 \end{pmatrix}$$

The original Matrix

# Example of Singular Value Decomposition

$$R_f = \begin{pmatrix} 1 & -1 & 1 & -1 & 1 & -1 \\ 1 & 1 & -0.2 & -1 & -1 & -1 \\ 0 & 1 & 1 & -1 & -1 & 0 \\ 1 & -1 & -1 & 1 & 1 & 1 \\ 1 & 0.2 & -1 & 1 & 1 & 1 \end{pmatrix}$$

Step 1: Fill missing values with the mean value of the row

# Example of Singular Value Decomposition

$$R_f = \begin{pmatrix} 1.0592 & -1.1604 & 0.9716 & -0.8515 & 0.8040 & -1.0592 \\ 0.6636 & 0.9039 & 0.5881 & -0.9242 & -1.1244 & -0.6636 \\ 0.4300 & 0.9623 & 0.3746 & -0.6891 & -1.1045 & -0.4300 \\ -0.9425 & -0.8181 & -0.8412 & 1.2010 & 1.1320 & 0.9425 \\ -1.0290 & -0.2095 & -0.9270 & 1.1475 & 0.5535 & 1.0290 \end{pmatrix}$$

Step 2: Apply SVD to the matrix

# Example of Singular Value Decomposition

$$R_f = \begin{pmatrix} 1 & -1 & 1 & -1 & 1 & -1 \\ 1 & 1 & 0.5881 & -1 & -1 & -1 \\ 0.4300 & 1 & 1 & -1 & -1 & -0.43 \\ 1 & -1 & -1 & 1 & 1 & 1 \\ 1 & -0.2095 & -1 & 1 & 1 & 1 \end{pmatrix}$$

Step 3. Modify the target matrix setting in the missing values the learn values and iterate **until convergence**




Figure 1: A mapping of movies based on two latent factors

# Factorization Models

**Fill the missing values with some values**

- Hard to do
- Inaccurate filling can distort the data

**Modeling directly the observed rating only??**



# Modeling directly the observed rating only??

- Stochastic gradient descent
  - For each given training set, the system predicts  $r_{ui}$ , and computes the associated prediction error

$$e_{ui} \stackrel{def}{=} r_{ui} - q_i^T p_u.$$

- Then it modifies the parameters by a magnitude proportional to  $\gamma$  in the opposite direction of the gradient, yielding:

$$\begin{aligned} q_i &\leftarrow q_i + \gamma \cdot (e_{ui} \cdot p_u - \lambda \cdot q_i) \\ p_u &\leftarrow p_u + \gamma \cdot (e_{ui} \cdot q_i - \lambda \cdot p_u) \end{aligned}$$

# Regularized SVD

- Direct Matrix Factorization of **Incomplete Data**
  - **Modeling directly the observed rating only**
  - When data is sparse, covariance estimates will be statistically unreliable.
- **Matrix factorization as a cost function**

$$\text{Min}_{p_u, q_i} \sum_{\text{known } r_{ui}} \left( r_{ui} - p_u^T q_i \right)^2 + \lambda \left( \|p_u\|^2 + \|q_i\|^2 \right)$$

- Optimize by either **stochastic gradient-descent** or **alternating least squares**
- The constant  **$\lambda$**  controls the extent of regularization

# Improved SVD

## Adding **Bias**

$$P'_{ui} = b_u + b_i + (p_u^T q_i)$$

$b_u$  is observed deviation of user  $u$  and  
 $b_i$  is observed deviation of item  $i$ .

Improving regularized singular value decomposition for collaborative filtering

A Paterek

Proceedings of KDD Cup and Workshop 2007, 5-8

713

2007

# SVD++

- Explicit + Implicit information
  - boolean implicit feedback, indicating if the item has been rated by the user or not.

$$\hat{r}_{ui} = b_{ui} + q_i^T \left( p_u + |N(u)|^{-\frac{1}{2}} \sum_{j \in N(u)} y_j \right), \text{ where } b_{ui} = \mu + b_u + b_i$$

- $N(u)$  is the set of items for which the user  $u$  has implicit information, and  $q$ ,  $p$  and  $y$  are latent factors

Factorization meets the neighborhood: a multifaceted collaborative filtering model

2008

2008

Y Koren

Proceedings of the 14th ACM SIGKDD international conference on Knowledge ...

# Non-Negative Matrix Factorization

- Can be used for ratings matrices that are non negative

$$\begin{array}{ll}\text{Minimize} & J = \frac{1}{2} \|R - UV^T\| \\ \text{subject to:} & U \geq 0, V \geq 0\end{array}$$

- The major advantage of this approach is not the accuracy, but that of the **high level of interpretability** it provides in understanding the user-item interactions.

# Matrix Factorization Methods

Method	Constraints	Objective	Advantages/Disadvantages
Unconstrained	No Constraints	Frobenius + regularizer	Highest quality solution Good for most matrices Regularization prevents overfitting Poor interpretability
SVD	Orthogonal Basis	Frobenius + regularizer	Good Visual Interpretability Out-of-sample recommendations Good for dense matrices Poor semantic interpretability Suboptimal in sparse matrices
Max Margin	No Constraints	Hinge Loss + margin regularizer	Highest quality solution Resists overfitting Similar to unconstrained Poor Interpretability Good for discrete ratings
NMF	Non Negativity	Frobenius + regularizer	Good quality solution High semantic interpretability Loses interpretability with like/dislike ratings Less overfitting in some cases Best for implicit feedback

# Sparse Linear Models (SLIM)

- Computes the item-item relations, by estimating an **item x item** sparse aggregation coefficient **matrix**  $S$ .
- The recommendation score of an unrated item  $i$  for a user  $u$  is:

$$\hat{r}_{ui} = \mathbf{r}_u^T \mathbf{S}_i.$$
$$\begin{array}{ll} \underset{S}{\text{minimize}} & \frac{1}{2} \sum_{u,i} (r_{ui} - \hat{r}_{ui})^2 + \frac{\beta}{2} \|S\|_F^2 + \lambda \|S\|_1, \\ \text{subject to} & S \geq 0, \text{ and} \\ & \text{diag}(S) = 0. \end{array}$$

**SLIM: Sparse linear methods for top-n recommender systems**

X Ning, G Karypis

Data Mining (ICDM), 2011 IEEE 11th International Conference on, 497-506

115

2011

# Sparse Linear Models (SLIM)

$$\hat{r}_i = \text{Matrix R} \times \text{Matrix S}$$

	t <sub>1</sub>	t <sub>2</sub>	t <sub>3</sub>	...	t <sub>N</sub>
u <sub>1</sub>	0	2	0	...	2
u <sub>2</sub>	?	3	2	...	0
...	...	...	...	...	2
u <sub>M</sub>	0	4	2	...	2

	t <sub>1</sub>	t <sub>2</sub>	...	t <sub>N</sub>
t <sub>1</sub>	0		...	
t <sub>2</sub>		0	...	
...	...	...	0	...
t <sub>N</sub>			...	0

$$\hat{r}_{ui} = \mathbf{r}_u^T \mathbf{s}_i.$$

$$\begin{aligned} & \underset{S}{\text{minimize}} && \frac{1}{2} \sum_{u,i} (r_{ui} - \hat{r}_{ui})^2 + \frac{\beta}{2} \|S\|_F^2 + \lambda \|S\|_1, \\ & \text{subject to} && S \geq 0, \text{ and} \\ & && \text{diag}(S) = 0. \end{aligned}$$



# Sparse Linear Models (SLIM)

- Although the SLIM method proposes a prediction model for the rating, the final use of the ratings is for **ranking** the items
- Generally used for data sets with unary ratings (clicks, buy,.. ) from implicit feedback
- Since, the weights are restricted to be positive, the impact of each weight on the score are highly interpretable

# Evaluation

# Evaluation

	User	Item	Rating
Training	U1	T1	4
	U1	T2	3
	U1	T3	3
	U2	T2	4
	U2	T3	5
	U2	T4	5
	U3	T4	5
Test	U1	T4	? (3)
	U2	T1	? (2)
	U3	T1	? (3)
	U3	T2	? (3.5)
	U3	T3	? (4)

$$P(U1, T4) = \text{Avg}(T4) = (5+4)/2 = 4.5$$

$$P(U2, T1) = \text{Avg}(T1) = 4/1 = 4$$

$$P(U3, T1) = \text{Avg}(T1) = 4/1 = 4$$

$$P(U3, T2) = \text{Avg}(T2) = (3+4)/2 = 3.5$$

$$P(U3, T3) = \text{Avg}(T3) = (3+5)/2 = 4$$

## 3. Evaluation by Metrics

$$\text{Mean Absolute Error (MAE)} = \frac{1}{n} \sum_{i=1}^n |e_i|$$

$$e_i = R(U, T) - P(U, T)$$

$$\text{MAE} = (|3 - 4.5| + |2 - 4| + |3 - 4| + |3 - 3.5| + |4 - 4|) / 5 = 1$$

# Rank Accuracy



- **Mean Average Precision (MAP)** calculates the precision at the position of every corrected item in the ranked results list

$$\text{AveP} = \frac{\sum_{k=1}^n (P(k) \times \text{rel}(k))}{\text{number of relevant documents}}$$

$$\text{MAP} = \frac{\sum_{q=1}^Q \text{AveP}(q)}{Q}$$

Query 1:   $\text{AP} = \frac{1}{3}(1/1 + 2/3 + 3/6) = 0.72$

Query 2:   $\text{AP} = \frac{1}{3}(1/1 + 2/2 + 3/4) = 0.917$

$$\text{MAP} = (0.72 + 0.917) / 2 = 0.8185$$

# Top-N Evaluation

Task: Top-N Items to a user U3

Training	User	Item	Rating
	U1	T1	4
	U1	T2	3
	U1	T3	3
	U2	T2	4
	U2	T3	5
	U2	T4	5
	U3	T4	5
Test	U1	T4	? (3)
	U2	T1	? (2)
	U3	T1	? (3)
	U3	T2	? (3.5)
	U3	T3	? (4)

$$P(U3, T1) = \text{Avg}(T1) = 4/1 = 4$$

$$P(U3, T2) = \text{Avg}(T2) = (3+4)/2 = 3.5$$

$$P(U3, T3) = \text{Avg}(T3) = (3+5)/2 = 4$$

**Real rank:** I3,I2,I1

**Predicted rank:** I3,I1,I2

Precision@N = #of hits/N

Precision@1 = 1/1

Precision@2 = 1/2

Precision@3 = 3/3

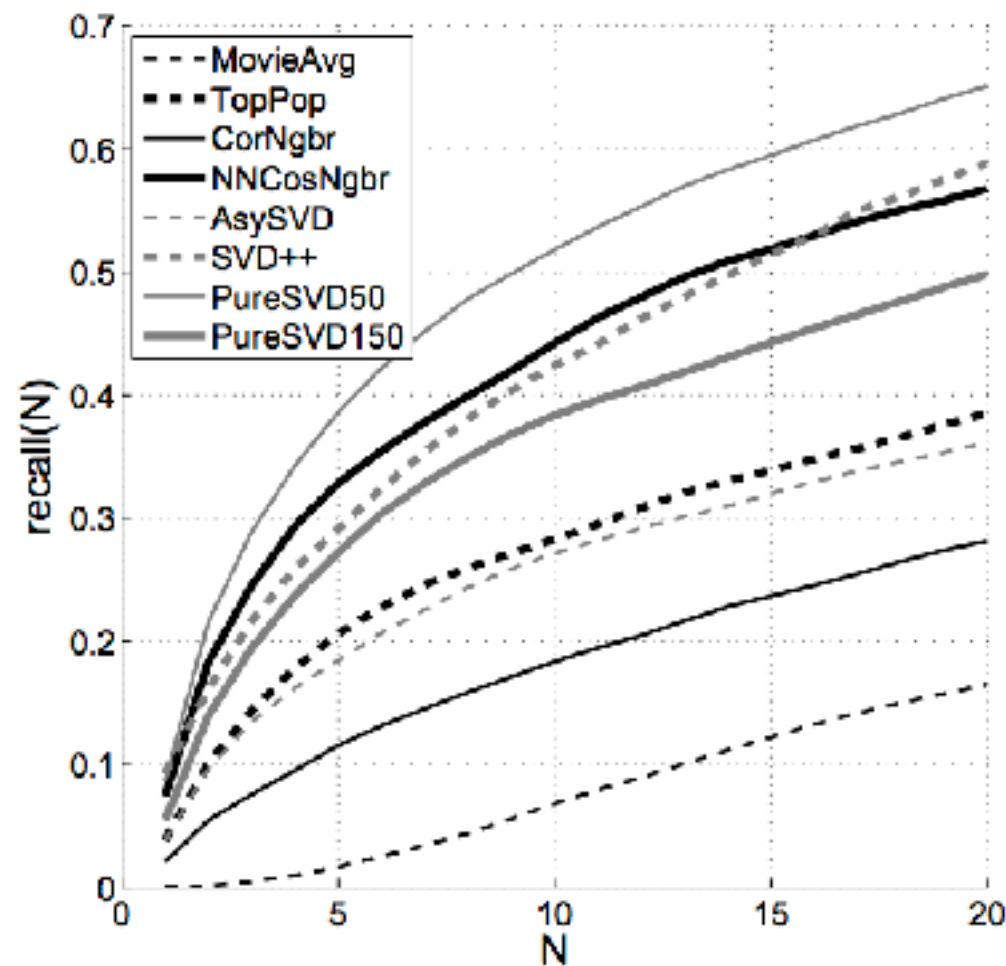
# Rank Accuracy




## Precision/Recall

- We can use the strategy followed by P. Cremonesi et.al.
- In order to measure the precision recall, first the models is trained using the training data, and then, for each item  $i$  rated with 5 stars in the test data set:
  - A set of 100 random unseen movies for the user of the item  $i$  are selected. We assume that these random movies will not be at the same interest than the 5 star movie
  - We predict the rating of the movie of item  $i$  and 100 random unseen movies.
  - We form a rank list by ordering all the 101 item according to the predicted rating. Let denote  $p$  the rank of the test item  $i$  within the list. The best results corresponds to the case the test item  $i$  precedes all the random items (i.e.,  $p=1$ ).
  - A top- $N$  recommendation list by piking the  $N$  top ranked items from the list. If  $p \leq N$  we have a hit. Otherwise we have a miss. Chances of hit increases as  $N$  is higher.

# Rank Accuracy



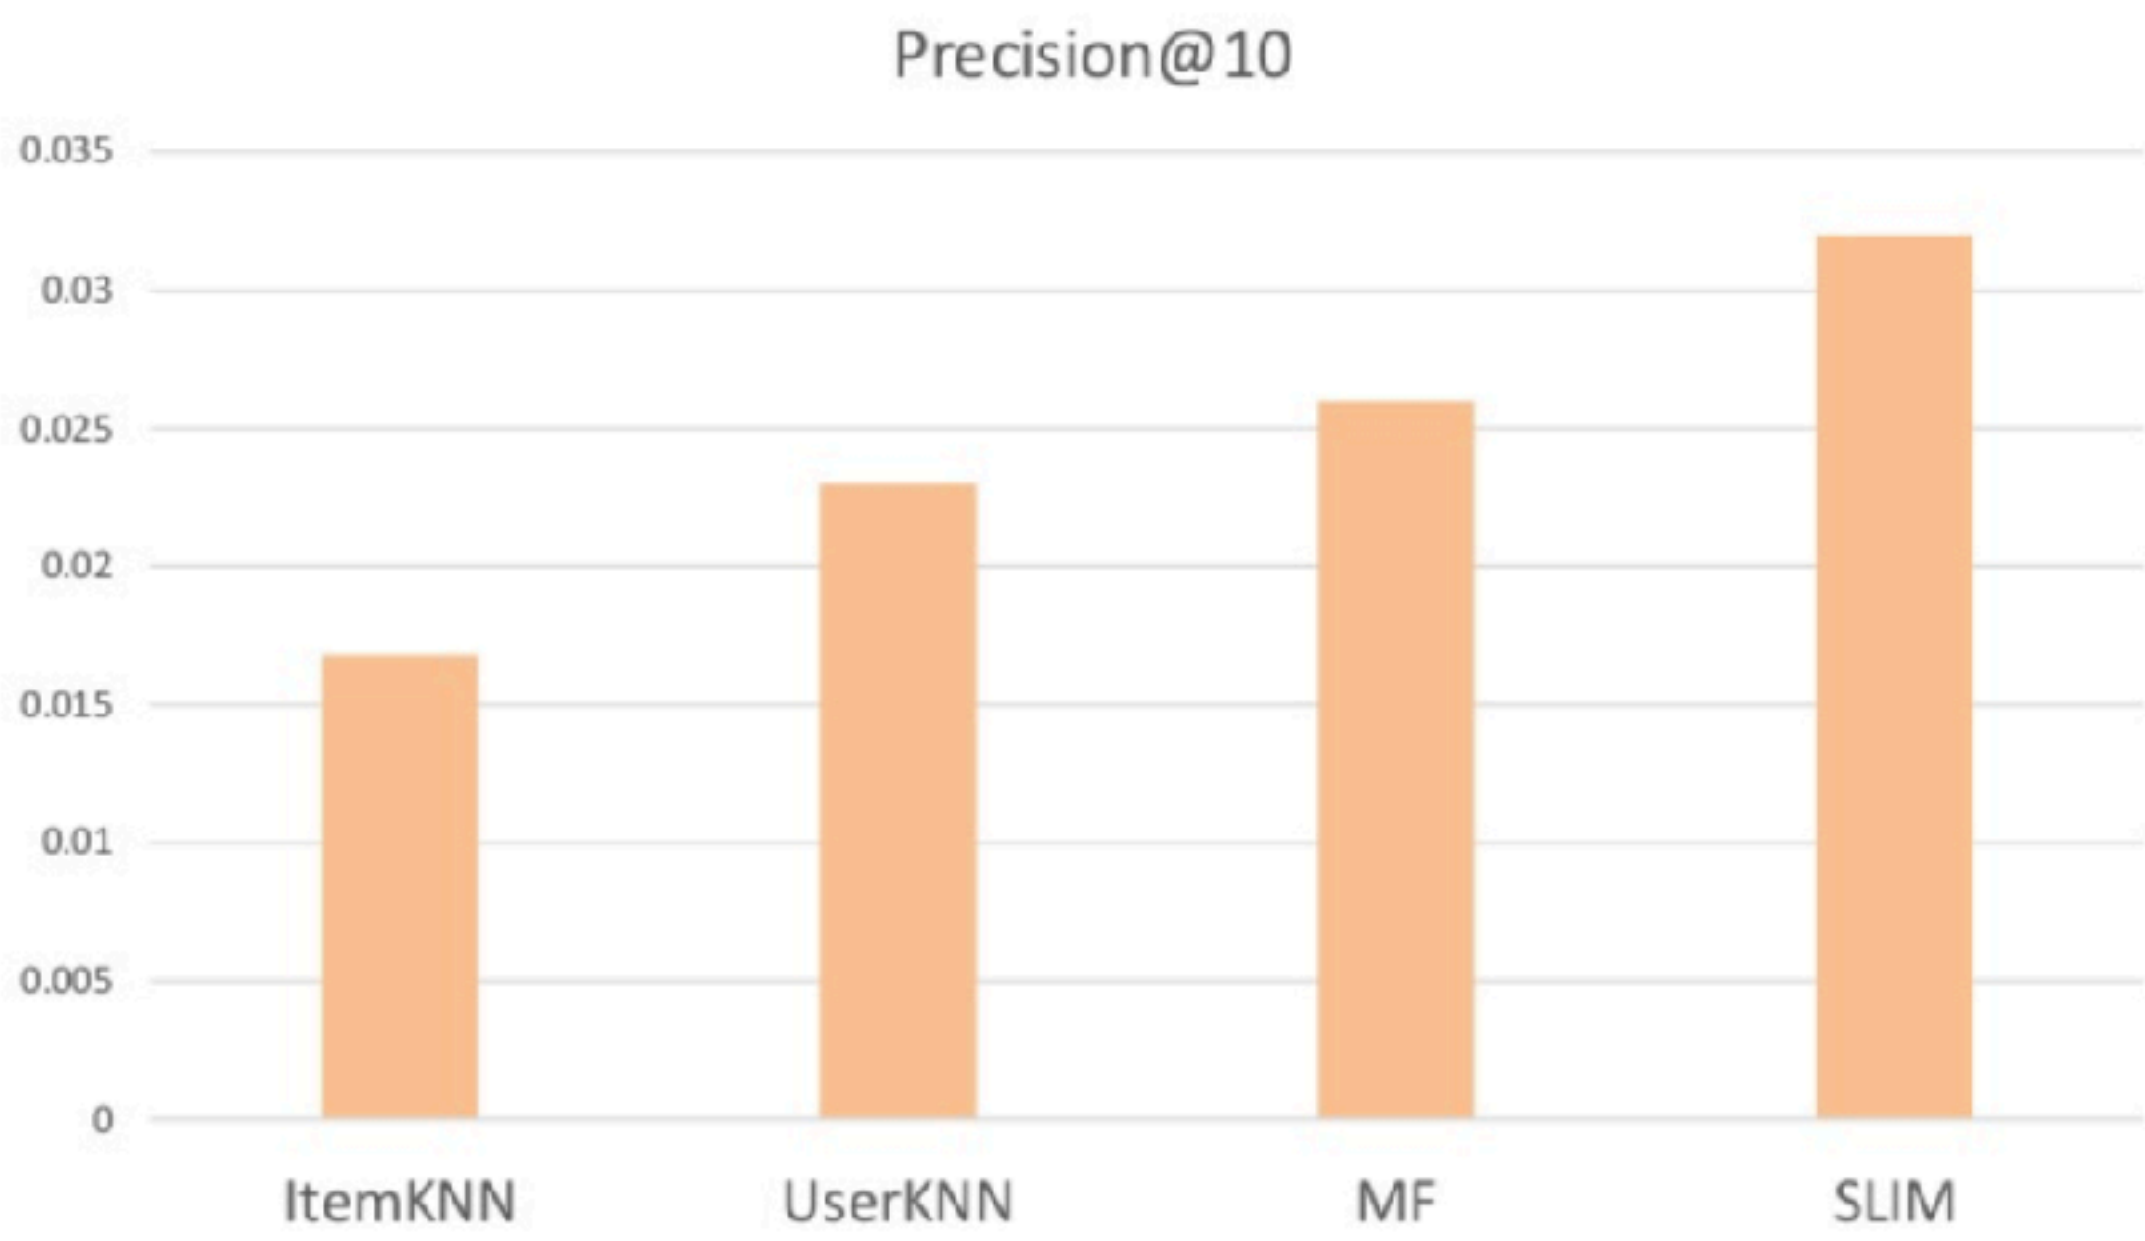
(a) recall



(b) precision vs recall

**Figure 2: Movielens: (a) recall-at- $N$  and (b) precision-versus-recall on all items.**

There are 100K ratings given by 943 users on 1,682 movies





# Task #2

## Movie Recommender

Create at least:

- 1) Item-Based Recommender System
- 2) Factorization Recommender system

### **What to deliver:**

Write a paper/report (4-8 pages)

Explain: 1) how you see the problem, 2) the used methods, 3) the results and 4) your conclusions in terms of accuracy and complexity

Github repository of the task.

Deadline - May 13th